





Register Allocation

Lecture 16

Prof. Aiken CS 143 Lecture 16

1

The Memory Hierarchy

	Register	0.3 ns	1-16 KB
	Cache L1	0.5 ns	8-64 KB
	Cache L2	7 ns	256-512 KB
	Cache L3	12 ns	1-8 MB
	Main memory	100 ns	1-32 GB
	HD seek	10M ns (10ms)	up to 10 TB
	SSD seek	0.1M ns (0.1ms)	up to 1 TB
	Network	150M ns (150 ms)	
			

CS 143 Lecture 16

4

Managing the Memory Hierarchy

- Most programs are written as if there are only two kinds of memory: main memory and disk
 - Programmer is responsible for moving data from disk to memory (e.g., file I/O)
 - Hardware is responsible for moving data between memory and caches
 - Compiler is responsible for moving data between memory and registers

Prof. Aiken CS 143 Lecture 16

6

Current Trends

- Power usage limits
 - Size and speed of registers/caches
 - Speed of processors
- But
 - The cost of a cache miss is very high
 - Typically requires 2-3 caches to bridge fast processor with large main memory
- It is very important to:
 - Manage registers properly
 - Manage caches properly
- Compilers are good at managing registers

Prof. Aiken CS 143 Lecture 16

7

The Register Allocation Problem

- Intermediate code uses unlimited temporaries
 - Simplifies code generation and optimization
 - Complicates final translation to assembly
- Typical intermediate code uses too many temporaries

Prof. Aiken CS 143 Lecture 16

8

The Register Allocation Problem (Cont.)

- The problem:
Rewrite the intermediate code to use no more temporaries than there are machine registers
- Method:
 - Assign multiple temporaries to each register
 - But without changing the program behavior

Prof. Aiken CS 143 Lecture 16

9

History

- Register allocation is as old as compilers
 - Register allocation was used in the original FORTRAN compiler in the '50s
 - Very crude algorithms
- A breakthrough came in 1980
 - Register allocation scheme based on graph coloring
 - Relatively simple, global and works well in practice

Prof. Aiken CS 143 Lecture 16

10

An Example

- Consider the program


```
a := c + d
e := a + b
f := e - 1
```

```
r1 := r2 + r3
r1 := r1 + r4
r1 := r1 - 1
```
- Can allocate **a**, **e**, and **f** all to one register (r_1):
 - Assume **a** and **e** dead after use
 - Temporary **a** can be "reused" after $e := a + b$
 - So can temporary **e**
 - A dead temporary is not needed
 - A dead temporary can be reused

Prof. Aiken CS 143 Lecture 16

11

The Idea

Temporaries t_1 and t_2 can share the same register if at any point in the program at most one of t_1 or t_2 is live.

Or

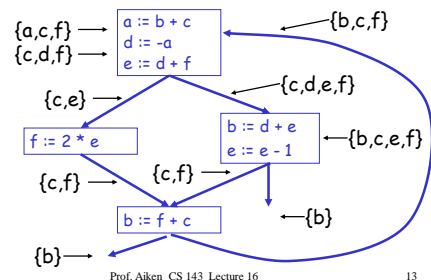
If t_1 and t_2 are live at the same time, they cannot share a register

Prof. Aiken CS 143 Lecture 16

12

Algorithm: Part I

- Compute live variables for each point:



Prof. Aiken CS 143 Lecture 16

13

The Register Interference Graph


- Construct an undirected graph
 - A node for each temporary
 - An edge between t_1 and t_2 if they are live simultaneously at some point in the program
- This is the *register interference graph* (RIG)
 - Two temporaries can be allocated to the same register if there is no edge connecting them

Prof. Aiken CS 143 Lecture 16

14

Example

- For our example:



- E.g., **b** and **c** cannot be in the same register
- E.g., **b** and **d** could be in the same register

Prof. Aiken CS 143 Lecture 16

15

Notes on Register Interference Graphs

- Extracts exactly the information needed to characterize legal register assignments
- Gives a global (i.e., over the entire flow graph) picture of the register requirements
- After RIG construction the register allocation algorithm is architecture independent

Prof. Aiken CS 143 Lecture 16

16

Definitions

- A coloring of a graph is an assignment of colors to nodes, such that nodes connected by an edge have different colors
- A graph is k-colorable if it has a coloring with k colors

Prof. Aiken CS 143 Lecture 16

17

Register Allocation Through Graph Coloring


- In our problem, colors = registers
 - We need to assign colors (registers) to graph nodes (temporaries)
- Let k = number of machine registers
- If the RIG is k -colorable then there is a register assignment that uses no more than k registers

Prof. Aiken CS 143 Lecture 16

18

Graph Coloring Example

- Consider the example RIG




- There is no coloring with less than 4 colors
- There are 4-colorings of this graph

Prof. Aiken CS 143 Lecture 16

19

Example Review




Prof. Aiken CS 143 Lecture 16

20

Example After Register Allocation

- Under this coloring the code becomes:



Prof. Aiken CS 143 Lecture 16

21

Computing Graph Colorings

- How do we compute graph colorings?
- It isn't easy:
 1. This problem is very hard (NP-hard). No efficient algorithms are known.
 - Solution: use heuristics
 2. A coloring might not exist for a given number of registers
 - Solution: later

Prof. Aiken CS 143 Lecture 16

22

Graph Coloring Heuristic

- Observation:
 - Pick a node t with fewer than k neighbors in RIG
 - Eliminate t and its edges from RIG
 - If resulting graph is k -colorable, then so is the original graph
- Why?
 - Let c_1, \dots, c_n be the colors assigned to the neighbors of t in the reduced graph
 - Since $n < k$ we can pick some color for t that is different from those of its neighbors

Prof. Aiken CS 143 Lecture 16

23

Graph Coloring Heuristic


- The following works well in practice:
 - Pick a node t with fewer than k neighbors
 - Put t on a stack and remove it from the RIG
 - Repeat until the graph has one node
- Assign colors to nodes on the stack
 - Start with the last node added
 - At each step pick a color different from those assigned to already colored neighbors

Prof. Aiken CS 143 Lecture 16

24

Graph Coloring Example (1)

- Start with the RIG and with $k = 4$:




- Remove a

Prof. Aiken CS 143 Lecture 16

25

Graph Coloring Example (2)




- Remove d

Prof. Aiken CS 143 Lecture 16

26

Graph Coloring Example (3)

- Note: all nodes now have fewer than 4 neighbors




- Remove c

Prof. Aiken CS 143 Lecture 16

27

Graph Coloring Example (4)



Stack: {c, d, a}

- Remove b

Prof. Aiken CS 143 Lecture 16

28

Graph Coloring Example (5)



Stack: {b, c, d, a}

- Remove e

Prof. Aiken CS 143 Lecture 16

29

Graph Coloring Example (6)



Stack: {e, b, c, d, a}

- Remove f

Prof. Aiken CS 143 Lecture 16

30

Graph Coloring Example (7)

- Now start assigning colors to nodes, starting with the top of the stack

Stack: {f, e, b, c, d, a}

Prof. Aiken CS 143 Lecture 16

31

Graph Coloring Example (8)



Stack: {e, b, c, d, a}

Prof. Aiken CS 143 Lecture 16

32

Graph Coloring Example (9)




Stack: {b, c, d, a}

- e must be in a different register from f

Prof. Aiken CS 143 Lecture 16

33


Graph Coloring Example (10)



Prof. Aiken CS 143 Lecture 16

34


Graph Coloring Example (11)



Prof. Aiken CS 143 Lecture 16

35

Graph Coloring Example (12)




- d can be in the same register as b

Prof. Aiken CS 143 Lecture 16

36

Graph Coloring Example (13)




Prof. Aiken CS 143 Lecture 16

37

What if the Heuristic Fails?

- What if all nodes have k or more neighbors ?
- Example: Try to find a 3-coloring of the RIG:




Prof. Aiken CS 143 Lecture 16

38

What if the Heuristic Fails?

- Remove a and get stuck (as shown below)
- Pick a node as a candidate for *spilling*
 - A spilled temporary "lives" in memory
 - Assume that f is picked as a candidate




Prof. Aiken CS 143 Lecture 16

39

What if the Heuristic Fails?

- Remove f and continue the simplification
 - Simplification now succeeds: b, d, e, c




Prof. Aiken CS 143 Lecture 16

40

What if the Heuristic Fails?

- Eventually we must assign a color to f
- We hope that among the 4 neighbors of f we use less than 3 colors \Rightarrow optimistic coloring



Prof. Aiken CS 143 Lecture 16

41

Spilling


- If optimistic coloring fails, we spill f
 - Allocate a memory location for f
 - Typically in the current stack frame
 - Call this address fa
- Before each operation that reads f , insert $f := \text{load } fa$
- After each operation that writes f , insert $\text{store } f, fa$

Prof. Aiken CS 143 Lecture 16

42

Spilling Example

- This is the new code after spilling f



Prof. Aiken CS 143 Lecture 16

43

A Problem


- This code reuses the register name f
- Correct, but suboptimal
 - Should use distinct register names whenever possible
 - Allows different uses to have different colors

Prof. Aiken CS 143 Lecture 16

44

Spilling Example

- This is the new code after spilling f

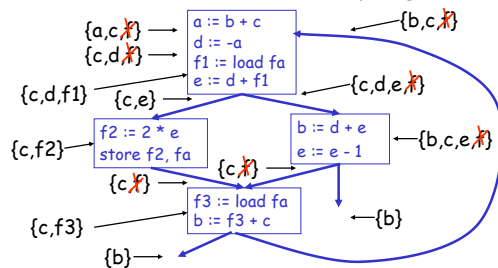


Prof. Aiken CS 143 Lecture 16

45

Recomputing Liveness Information

- The new liveness information after spilling:



Prof. Aiken CS 143 Lecture 16

46

Recomputing Liveness Information


- New liveness information is almost as before
 - Note f has been split into three temporaries
- $f1$ is live only
 - Between a $f1 := \text{load } fa$ and the next instruction
 - Between a $\text{store } f1, fa$ and the preceding instr.
- Spilling reduces the live range of f
 - And thus reduces its interferences
 - Which results in fewer RIG neighbors

Prof. Aiken CS 143 Lecture 16

47

Recompute RIG After Spilling

- Some edges of the spilled node are removed
- In our case f still interferes only with c and d
- And the resulting RIG is 3-colorable



Prof. Aiken CS 143 Lecture 16

48

Spilling Notes

- Additional spills might be required before a coloring is found
- The tricky part is deciding what to spill
 - But any choice is correct
- Possible heuristics:
 - Spill temporaries with most conflicts
 - Spill temporaries with few definitions and uses
 - Avoid spilling in inner loops

Prof. Aiken CS 143 Lecture 16

49

Caches

- Compilers are very good at managing registers
 - Much better than a programmer could be
- Compilers are not good at managing caches
 - This problem is still left to programmers
 - It is still an open question how much a compiler can do to improve cache performance
- Compilers can, and a few do, perform some cache optimizations

Prof. Aiken CS 143 Lecture 16

50

Cache Optimization

- Consider the loop


```
for(j := 1; j < 10; j++)
  for(i=1; i<1000000; i++)
    a[i] *= b[i]
```
- This program has terrible cache performance
 - Why?

Prof. Aiken CS 143 Lecture 16

51

Cache Optimization (Cont.)

- Consider the program:

```
for(i:=1; i<1000000; i++)  
  for(j := 1; j < 10; j++)  
    a[i] *= b[j]
```

 - Computes the same thing
 - But with much better cache behavior
 - Might actually be more than 10x faster
- A compiler can perform this optimization
 - called loop interchange

Conclusions

- Register allocation is a “must have” in compilers:
 - Because intermediate code uses too many temporaries
 - Because it makes a big difference in performance
- Register allocation is more complicated for CISC machines