

Language Security

Lecture 18

Prof. Aiken CS 143 Lecture 18

1

Lecture Outline

- Beyond compilers
 - Looking at other issues in programming language design and tools
- C
 - Arrays
 - Exploiting buffer overruns
 - Detecting buffer overruns

Prof. Aiken CS 143 Lecture 18

2

Platitudes

- Language design has influence on
 - Safety
 - Efficiency
 - Security

Prof. Aiken CS 143 Lecture 18

3

C Design Principles

- Small language
- Maximum efficiency
- Safety less important
- Designed for the world in 1972
 - Weak machines
 - Trusted networks


Prof. Aiken CS 143 Lecture 18

4

Arrays in C

```
char buffer[100];
```

Declares and allocates an array of 100 chars



Prof. Aiken CS 143 Lecture 18

5

C Array Operations

```
char buf1[100], buf2[100];
```

Write:

```
buf1[0] = 'a';
```

Read:

```
return buf2[0];
```

Prof. Aiken CS 143 Lecture 18

6

What's Wrong with this Picture?

```
int i = 0;
for(i = 0; buf1[i] != '\0'; i++)
    { buf2[i] = buf1[i]; }
buf2[i] = '\0';
```

Prof. Aiken CS 143 Lecture 18

7

Indexing Out of Bounds

The following are all legal C and may generate no run-time errors

```
char buffer[100];

buffer[-1] = 'a';
buffer[100] = 'a';
buffer[100000] = 'a';
```

Prof. Aiken CS 143 Lecture 18

8

Why?

- Why does C allow out of bounds array references?
 - Proving at compile-time that all array references are in bounds is very difficult (impossible in C)
 - Checking at run-time that all array references are in bounds is expensive

Prof. Aiken CS 143 Lecture 18

9

Code Generation for Arrays

```
buf1[i] = 1; /* buf1 has type int[] */

r1 = load &buf1;
r2 = load i;
r3 = r2 * 4;
r4 = r1 + r3
store r4, 1
```

Prof. Aiken CS 143 Lecture 18

10

Discussion

- 5 instructions worst case
- Often `&buf1` and `i` already in registers
 - Saves 2 instructions
- Many machines have indirect loads/stores
 - `store r1[r3], 1`
 - Saves 1 instruction
- Best case 2 instructions
 - Offset calculation and memory operation

Prof. Aiken CS 143 Lecture 18

11

Code Generation for Arrays with Bounds Checks

```
buf1[i] = 1; /* buf1 has type int[] */

r1 = load &buf1;
r2 = load i;
r3 = r2 * 4;
if r3 < 0 then error;
r5 = load limit of buf1;
if r3 >= r5 then error;
r4 = r1 + r3
store r4, 1
```

Prof. Aiken CS 143 Lecture 18

12

Discussion

- Lower bounds check can often be removed
 - Easy to prove statically that index is positive
- Upper bounds check hard to remove
 - Leaves a conditional in instruction stream
- In C, array limits not stored with array
 - Knowing the array limit for a given reference is non-trivial

Prof. Aiken CS 143 Lecture 18

13

C vs. Java

- C array reference typical case
 - Offset calculation
 - Memory operation (load or store)
- Java array reference typical case
 - Offset calculation
 - Memory operation (load or store)
 - Array bounds check
 - Type compatibility check (for stores)

Prof. Aiken CS 143 Lecture 18

14

Buffer Overruns

- A buffer overrun writes past the end of an array
- *Buffer* usually refers to a C array of char
 - But can be any array
- So who's afraid of a buffer overrun?
 - Can damage data structures
 - Cause a core dump
 - What else?

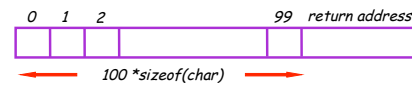
Prof. Aiken CS 143 Lecture 18

15

Stack Smashing

Buffer overruns can alter the control flow of your program!

`char buffer[100]; /* stack allocated array */`



Prof. Aiken CS 143 Lecture 18

16

An Overrun Vulnerability


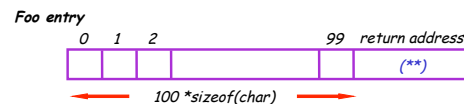
```
void foo(char buf1[]) {
    char buf2[100];
    int i = 0;
    for(i = 0; buf1[i] != '\0'; i++)
        { buf2[i] = buf1[i]; }
    buf2[i] = '\0';
}
```

Prof. Aiken CS 143 Lecture 18

17

An Interesting Idea

```
char buf[104] = { '\0', ..., '\0', magic 4 chars }
foo(buf); (**)
```



Prof. Aiken CS 143 Lecture 18

18

Discussion

- So we can make `foo` jump wherever we like.
- How is this possible?
- Unanticipated interaction of two features:
 - Unchecked array operations
 - Stack-allocated arrays
 - Knowledge of frame layout allows prediction of where array and return address are stored
 - Note the “magic cast” from char’s to an address

Prof. Aiken CS 143 Lecture 18

19

The Rest of the Story

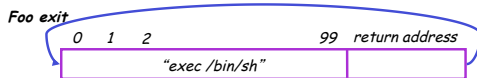
- We can make `foo` jump anywhere.
- But where is a useful place to jump?
- Idea: Put our own code in the buffer and jump there!

Prof. Aiken CS 143 Lecture 18

20

The Plan

```
char buf[104] = { 104 magic chars }  
foo(buf);
```



Prof. Aiken CS 143 Lecture 18

21

Details

- “`exec /bin/sh`”
 - Easy to write in assembly code
 - Make all jumps relative
- Be careful not to have null's in the code (why?)

Prof. Aiken CS 143 Lecture 18

22

More Details

- Overwrite return address with start of buffer
 - Harder
 - Need to guess where buffer in called routine starts (trial & error)
 - Pad front of buffer with NOPS
 - Guess need not be exact; just land somewhere in NOPS

Prof. Aiken CS 143 Lecture 18

23

And More Details

- Overwrite return address
 - Don't need to know exactly where return address is
 - Just pad end of buffer with multiple copies of new return address X

```
char buf[104] =  
    "NOPS ... /bin/exec sh XXXXXXXXXXXX"  
foo(buf);
```

Prof. Aiken CS 143 Lecture 18

24

The State of C Programming

- Buffer overruns are common
 - Programmers must do their own bounds checking
 - Easy to forget or be off-by-one or more
 - Program still appears to work correctly
- In C wrt to buffer overruns
 - Easy to do the wrong thing
 - Hard to do the right thing

Prof. Aiken CS 143 Lecture 18

25

The State of Hacking

- Buffer overruns widely known since the 1980's
 - Remain a popular attack today
- Highly automated toolkits available to exploit known buffer overruns
 - Search for "buffer overruns" yields > 100,000 hits

Prof. Aiken CS 143 Lecture 18

26

The Sad Reality

- Even well-known buffer overruns are still exploited
 - Hard to get people to upgrade millions of vulnerable machines
- We assume that there are many more unknown buffer overrun vulnerabilities
 - At least unknown to the good guys

Prof. Aiken CS 143 Lecture 18

27

How Do We Prevent Buffer Overruns?

- Many proposed techniques!
 - A research Rorschach test
- A brief survey
 - Language design
 - Static analysis
 - Dynamic analysis

Prof. Aiken CS 143 Lecture 18

28

Language Design

- Enforce data abstractions!
- How?
 - Type safety
 - The guarantee that if $e:T$, then e evaluates to a value of type T
 - No unsafe casts
 - Memory safety
 - Array bounds checking
 - No computation on pointers
 - Automatic memory management

Prof. Aiken CS 143 Lecture 18

29

Tools for Static Memory Safety

- Bug finding tools
 - Detect common patterns of buffer overruns
 - Use heuristics
 - Focus on scenarios likely to be real overruns, rather than obscure scenarios that might not be
 - Avoid *false positives*
- Verification
 - Formally prove memory safety
 - Can require deep understanding of the program's semantics

Prof. Aiken CS 143 Lecture 18

30

Dynamic Memory Safety

- Many proposals
- Sandboxing
 - Confine all memory references in the program to its own data space
 - Guarantees damage is limited to the program itself
- Code and data randomization
 - Give everyone a slightly different binary and data layout
 - Variation minimizes chances an attack can work on all copies of a program

Prof. Aiken CS 143 Lecture 18

31

Summary

- Programming language knowledge useful beyond compilers
- Useful for programmers
 - Understand what you are doing!
- Useful for tools other than compilers
 - Big research direction

Prof. Aiken CS 143 Lecture 18

32